

Query Processing and Optimization in SQL

Rishit Garg	Nived Shah	Harshit Jain	Sachish Singla	Parth Patil
22CS30045	22CS10049	22CS10030	22CS30046	22CS30041

Team PRQL

Abstract

As the scale of data in modern applications continues to grow, inefficient SQL queries can become significant bottlenecks in database performance. This project aims to implement a custom SQL query optimization system that enhances query execution by integrating parsing, relational algebra transformations, and cost-based optimization strategies. The system is designed to demonstrate practical applications of classical database optimization techniques.

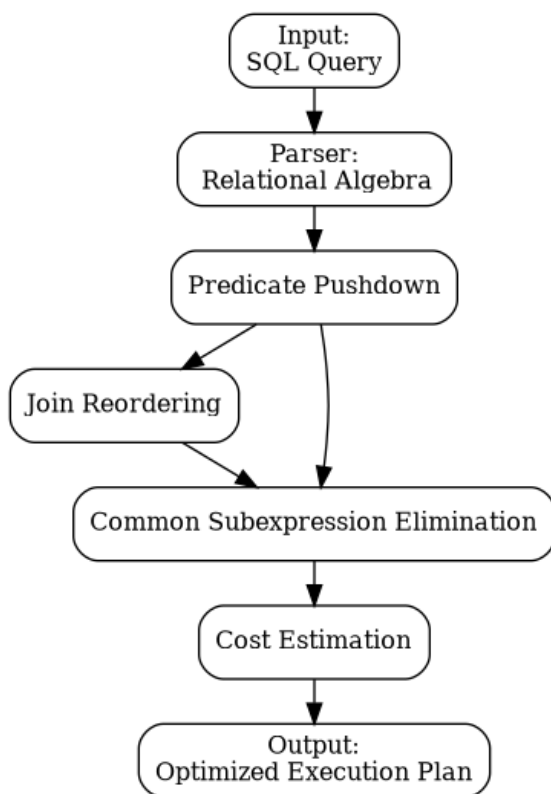


Figure 1: Query Optimization Pipeline

Methodology

This system was architected with a modular, layered approach to enable effective, step-wise SQL query optimization. Each component focuses on a distinct optimization technique, working together to transform raw SQL input into an optimized relational algebra plan. **PostgreSQL** plays a central role in this pipeline as a source of crucial statistical metadata used in **cost-based decisions**. The end-to-end pipeline begins with SQL parsing and proceeds through logical transformations such as predicate push-down and common subexpression elimination, culminating in a cost-based join optimization step. At each stage, intermediate plans are visualized using a web interface. For an overview of the full pipeline, refer to Figure 1.

1. SQL Parser

The SQL parser is the first component of the query optimization pipeline. Its main function is to tokenize SQL queries and transform them into a machine-readable structure: a **relational algebra tree** represented in **JSON format**. This structured output forms the basis for all subsequent optimizations.

Tools Used:

- **Flex**: Tokenizes the raw SQL input by breaking it down into meaningful units such as keywords (SELECT, FROM, WHERE), identifiers (table/column names), operators, and literals.
- **Bison**: Uses a context-free grammar to recognize SQL syntax and build an abstract syntax tree (AST). The AST is then transformed into a relational algebra JSON.

Supported SQL Subset

The parser currently supports a meaningful subset of SQL that includes:

- Basic SELECT-FROM-WHERE queries
- Simple JOINS (using ON operator)
- Nested subqueries

Example SQL Query Input and Output

```
SELECT L.L_ORDERKEY, L.L_PARTKEY
FROM LINEITEM L
WHERE L.L_SHIPDATE >= '1994-01-01'
```

```
{
  "condition": {
    "left": {
      "attr": "L_SHIPDATE",
      "table": "L"
    },
    "right": {
      "type": "string",
      "value": "1994-01-01"
    },
    "type": "GE"
  },
  "input": {
    "columns": [
      {
        "attr": "L_ORDERKEY",
        "table": "L"
      },
      {
        "attr": "L_PARTKEY",
        "table": "L"
      }
    ],
    "input": {
      "tables": [
        {
          "alias": "L",
          "name": "LINEITEM"
        }
      ],
      "type": "table_scan"
    },
    "type": "project"
  },
  "type": "select"
}
```

2. Predicate Pushdown

The second phase of the query optimization pipeline focuses on a classical and highly effective transformation in relational algebra: **Predicate Pushdown**. This optimization rule aims to minimize the size of the data retrieved from the storage system. The rule reduces the size of data that are fetched by applying the filter (predicate) as near as possible to table scanning. This additionally reduces the size of intermediate results that optimize other operations such as JOIN.

Motivation

In the original un-optimized query tree generated by the SQL parser, **Select** operations may appear higher up in the tree, often after expensive operations such as joins. Today, there is an increasing use of cloud-based and remote database systems, which calls for a method to reduce **network transfer latency** and data retrieval process (to provide a **higher degree of concurrency**). This is why Predicate Pushdown is an effective approach since it reduces the amount of data is fetched and transferred across the network from the remote host to the client. By moving down the selection condition closer to the relevant base tables so, the system:

- Filters irrelevant rows early
- Shrinks the size of data being passed through joins or projections
- Reduces memory consumption and improves processing speed

Algorithm 1 Predicate Pushdown Optimization

Input: Logical query plan tree T

Output: Optimized query plan with predicates pushed down

function PUSHDOWN(T)

if T is a leaf node (e.g., TableScan) **then**

return T

end if

if T is a Select node with condition C over subtree S **then**

 Decompose C into conjunctive components $C_1 \wedge C_2 \wedge \dots \wedge C_n$

$S \leftarrow \text{PUSHDOWN}(S)$

for all C_i in $\{C_1, C_2, \dots, C_n\}$ **do**

 Identify the minimal subtree S_i that contains all referenced attributes in C_i

 Push C_i down to just above S_i

end for

return Updated subtree with conditions pushed

else

for all child node T_c of T **do**

$T_c \leftarrow \text{PUSHDOWN}(T_c)$

end for

return Updated node with pushed-down children

end if

end function

Example Query:

```
SELECT L.L_ORDERKEY, L.L_PARTKEY
FROM LINEITEM L
JOIN ORDERS O ON L.L_ORDERKEY = O.O_ORDERKEY
WHERE L.L_SHIPDATE >= '1994-01-01' AND O.O_ORDERSTATUS = 'F'
```

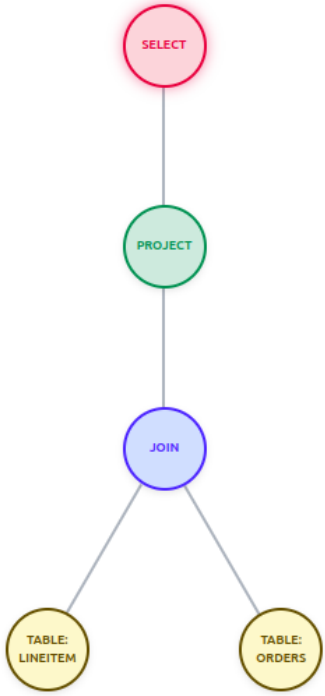


Figure 2: Original Query Plan

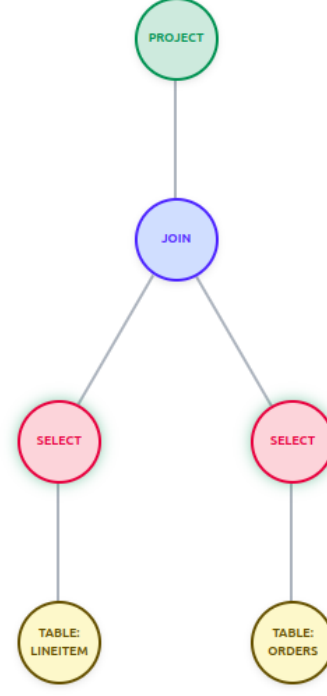


Figure 3: After Predicate Pushdown

For a detailed explanation of the original and optimized query plans, refer to the appendix.

3. Join Optimizer

The join optimizer is the most computation-intensive phase in the query optimization pipeline. It takes as input the optimized relational algebra tree (post predicate pushdown) in JSON format and outputs a reordered join plan that minimizes query execution cost using cost-based optimization.

Motivation

Join operations are among the most expensive in SQL query processing. The order in which joins are executed can have a drastic impact on performance, especially for queries involving multiple tables. An inefficient join ordering may produce large intermediate results, whereas an optimal ordering can significantly reduce computation.

Join Predicate Graph Construction

The optimizer first constructs a **join predicate graph**. This graph-based representation is crucial for understanding all relations involved and their predicate dependencies, including **implicit transitive joins** that might not be obvious from the original SQL. To demonstrate, here is an example set of join predicates:

```

A JOIN B ON A.id = B.id2
B JOIN C ON B.id2 = C.id3
  
```

From the above, the optimizer infers that:

- A joins with B (using `id2` of B)
- B joins with C (using `id2` of B)
- Therefore, A can indirectly join with C via B (transitive join)

This transitive inference allows the exploration of additional valid join paths (e.g., A-C) that may yield lower costs when reordered.

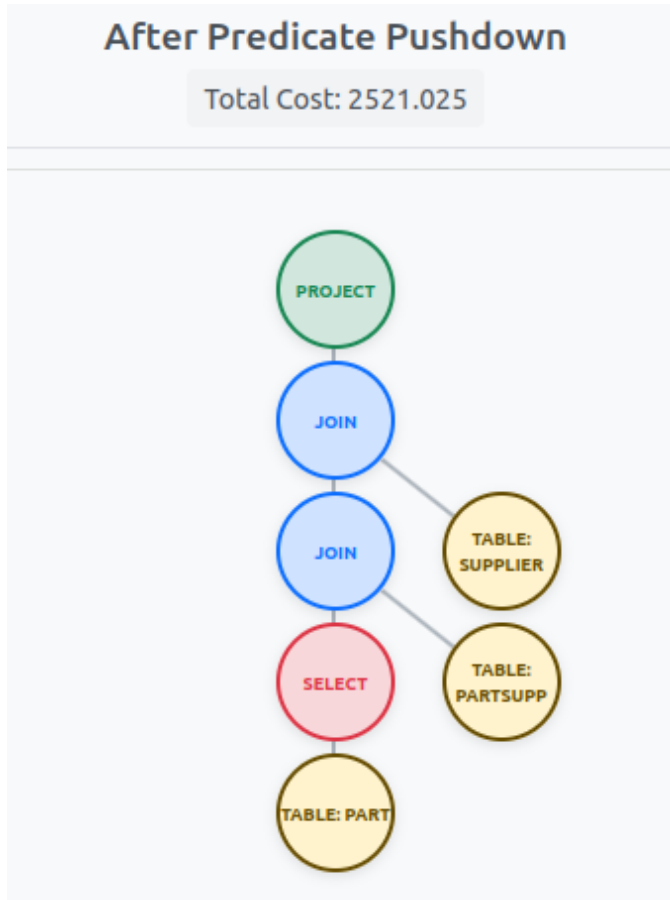


Figure 4: Original Query Plan



Figure 5: After Join Reordering

Join Order Enumeration and Strategy

Once the join graph is built, the optimizer systematically explores all possible join orderings. For each valid order (based on tables permutations), it evaluates different join strategies and uses PostgreSQL's statistics (from `'pg_stat'`) to estimate cost.

Steps:

1. **Cost Estimator Models:** Iterate through different models for cost estimation:
 - **MCV (Most Common Value)** – focuses on skewed distributions

- **NDV (Number of Distinct Values)** – emphasizes cardinality spread
 - **Fixed Estimates** – fallback with predefined selectivity assumptions
2. **Join Orderings:** Generate all valid permutations of table join sequences (e.g., ABC, ACB, BAC, etc.)
 3. **Join Tree Traversal:** For each order, simulate **left-deep join trees** using dynamic programming (Selinger-style):
 - At each step, form a binary join node ($R \bowtie S$)
 - Use **PostgreSQL statistics** for R and S to calculate the node cost.
 4. **Join Strategies:** For each binary join:
 - **Hash Join** – best when hashable keys exist
 - **Nested Loop Join** – efficient for small tables
 - **Block Nested Loop** – a middle ground strategy
 5. **Cost Estimation:** Compute cost as a combination of:
 - Estimated cardinality
 - Disk I/O
 - CPU usage
 - Selectivity of predicates
 6. **Dynamic Programming (Selinger’s Algorithm):**
 - Maintain a cost table mapping subplans to their best-known cost
 - Reuse optimal subtrees to avoid recomputation

Algorithm 2 Selinger-style Dynamic Programming Join Order Optimization

```

Initialize cost table  $C$ 
for each single relation  $R$  do
  Set  $C[\{R\}] = \text{cost}(R)$ 
end for
for each join size  $s = 2$  to  $n$  do
  for each subset  $S$  of  $s$  relations do
    for each partition of  $S$  into  $S_1$  and  $S_2$  do
      Compute cost of joining  $C[S_1]$  and  $C[S_2]$ 
      Update  $C[S]$  if this plan is cheaper
    end for
  end for
end for
Return optimal plan from  $C[\text{all relations}]$ 
  
```

4. Common Subexpression Elimination

In the context of relational algebra and SQL queries, common subexpressions typically arise when the same logical operations—such as selections, projections, or joins—are applied multiple times across different parts of a query.

Motivation

Common Subexpression Elimination not only simplifies the logical representation of the query but also facilitates more efficient physical execution by allowing intermediate results to be computed once and reused.

Redundancy Detection

The system traverses the relational algebra tree to identify structurally and semantically identical subtrees—i.e., common subexpressions. These subtrees are compared based on their operation types, inputs, and conditions, while trivial variations (e.g., different column names or aliases) are ignored. Only subtrees of sufficient depth and recurrence are considered meaningful for elimination. This process ensures that only computationally expensive or repeatedly occurring expressions are selected for optimization.

Replacement by References and DAG Construction

Once redundancies are detected, each unique subexpression is abstracted and assigned a symbolic identifier. Every instance of the same expression in the tree is then replaced with a reference to this identifier. To support such reference-based reuse, the original tree structure is transformed into a Directed Acyclic Graph (DAG) following the steps below:

- The tree is traversed in a bottom-up fashion.
- A structural hash is computed for each subtree to detect equivalence.
- When a match is found, the existing node is reused instead of duplicating it.

This DAG representation allows the system to compute shared subexpressions only once, reducing redundancy, saving computation, and improving query execution efficiency without altering semantics.

Example Query:

```
SELECT H.C_NAME, O.O_ORDERKEY, O.O_TOTALPRICE
FROM (
  SELECT CUSTOMER.C_CUSTKEY, CUSTOMER.C_NAME, CUSTOMER.C_ACCTBAL
  FROM CUSTOMER
  WHERE CUSTOMER.C_ACCTBAL > 1000
) H
JOIN (
  SELECT CUSTOMER.C_CUSTKEY, CUSTOMER.C_NAME, CUSTOMER.C_ACCTBAL
  FROM CUSTOMER
  WHERE CUSTOMER.C_ACCTBAL > 1000
```



```

) I ON H.CUSTOMER.C_NATIONKEY = I.CUSTOMER.C_NATIONKEY
JOIN ORDERS O ON H.CUSTOMER.C_CUSTKEY = O.O_CUSTKEY
WHERE O.O_TOTALPRICE > 50000 AND O.O_ORDERSTATUS = 'O' OR O.O_SHIPPRIORITY = 0

```

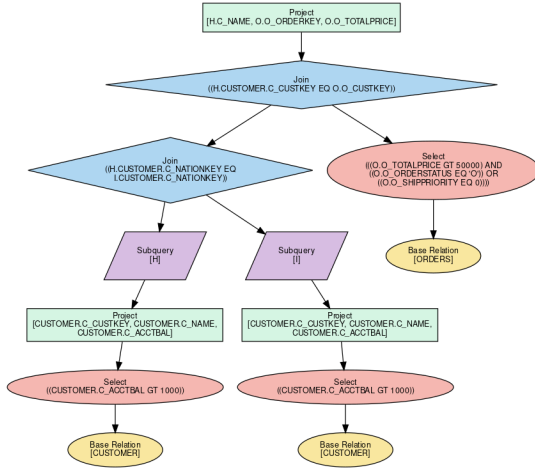


Figure 6: After Predicate Pushdown

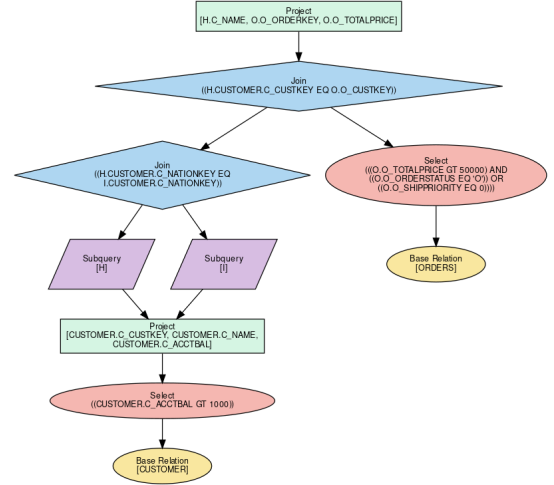


Figure 7: After CSE

5. Web Interface for Visualization

To facilitate user comprehension and interaction, an intuitive web interface was designed to visualize each stage of the SQL query optimization pipeline. This interface dynamically renders relational algebra trees corresponding to different phases: parsing, predicate pushdown, join reordering, and common subsequence elimination.

Technologies Used

- **Flask:** Serves as the backend framework, handling HTTP requests and responses, and orchestrating the flow between different optimization stages.
- **Graphviz/JavaScript:** Generate visual representations of the relational algebra trees.
- **HTML/CSS:** Construct the frontend, providing a responsive and interactive user interface.

Testing Methodology and Results

To test the robustness and effectiveness of our Query Optimization system, we chose to use the schema provided in the **TPC-H Benchmark**. The schema has been provided in the appendix for reference. Since our system runs on a specific subset of queries, we chose to create our own set of 10 queries aimed at rigorously testing every aspect of the system. Overall, we achieved an average of **18% reduction in cost**.

Conclusion and Future Scope

We have implemented a modular query optimization pipeline that processes SQL into relational algebra and applies classical and cost-based transformations. Each stage (parsing, pushdown, join ordering and common subexpression elimination) outputs JSON, facilitating debugging and extensibility. The system was tested on the TPC-H Benchmark for validation. The web interface built with Flask allows intuitive visualization of the entire optimization process.

Future Improvements:

- **Support for additional SQL constructs (e.g., GROUP BY, WITH):** Such additional constructs give more variety of costing strategy that can be used to annotate the query plans.
- **Increasing query plan search space:** With the current approach, we find all possible Join Orders and evaluate the one with the least cost. However, cost-based optimization also includes generating all possible query plans considering other operators too, and finding the least cost plan amongst them.
- **Utilizing Machine Learning techniques for choosing cost parameters:** With knowledge of access history, the cost parameters such as predicate-selectivity can be predicted dynamically based on input query. This can be an interesting field to understand speed vs cost-prediction tradeoff since query fetching time must be fast, but including dynamicity for calculation may reduce speed.
- **Understanding physical storage and Indexing support:** By default, Postgres can perform indexing based on primary-key. Identifying queries with use of primary-keys can enable use of Indexing to reduce cost. Additionally, strategies can be explored to create Indexes for high-volume data thus assisting clients to focus on abstract querying. This requires a thorough understanding of the physical storage used and data structures supported, and evaluating trade-offs between index creation cost and query processing cost dynamically.

References

1. Silberschatz, Korth, and Sudarshan. *Database System Concepts*. 7th Edition.
2. Jan Van den Bussche, Stijn Vansummeren. *Translating SQL into the Relational Algebra*
3. PostgreSQL Documentation: <https://www.postgresql.org/docs/>
4. TPC-H Documentation: <https://www.tpc.org/tpch/default5.asp>
5. tcph dbgen: <https://github.com/electrum/tpch-dbgen>
6. Flask Documentation: <https://flask.palletsprojects.com/en/stable/>
7. Graphviz Documentation: <https://graphviz.org/documentation/>

Appendix

Explanation of Predicate Pushdown Example

Original Plan

```
FILTER [(L.L_SHIPDATE >= '1994-01-01') AND (O.O_ORDERSTATUS = 'F')]
  PROJECT ['L.L_ORDERKEY', 'L.L_PARTKEY']
    JOIN [L.L_ORDERKEY = O.O_ORDERKEY]
      SCAN(LINEITEM AS L)
      SCAN(ORDERS AS O)
```

Optimized Plan

```
PROJECT ['L.L_ORDERKEY', 'L.L_PARTKEY']
  JOIN [L.L_ORDERKEY = O.O_ORDERKEY]
    FILTER [L.L_SHIPDATE >= '1994-01-01']
      SCAN(LINEITEM AS L)
    FILTER [O.O_ORDERSTATUS = 'F']
      SCAN(ORDERS AS O)
```

TPC-H Schema

```
CREATE TABLE NATION (
  N_NATIONKEY  INTEGER      NOT NULL,
  N_NAME       CHAR(25)     NOT NULL,
  N_REGIONKEY  INTEGER      NOT NULL,
  N_COMMENT    VARCHAR(152)
);
```

```
CREATE TABLE REGION (
  R_REGIONKEY  INTEGER      NOT NULL,
  R_NAME       CHAR(25)     NOT NULL,
  R_COMMENT    VARCHAR(152)
);
```

```
CREATE TABLE PART (
  P_PARTKEY    INTEGER      NOT NULL,
  P_NAME       VARCHAR(55)   NOT NULL,
  P_MFGR       CHAR(25)     NOT NULL,
  P_BRAND      CHAR(10)     NOT NULL,
  P_TYPE       VARCHAR(25)   NOT NULL,
  P_SIZE       INTEGER      NOT NULL,
  P_CONTAINER  CHAR(10)     NOT NULL,
  P_RETAILPRICE DECIMAL(15,2) NOT NULL,
```

```
        P_COMMENT      VARCHAR(23)      NOT NULL
    );
```

```
CREATE TABLE SUPPLIER (
    S_SUPPKEY      INTEGER      NOT NULL,
    S_NAME         CHAR(25)     NOT NULL,
    S_ADDRESS      VARCHAR(40)  NOT NULL,
    S_NATIONKEY    INTEGER      NOT NULL,
    S_PHONE        CHAR(15)     NOT NULL,
    S_ACCTBAL      DECIMAL(15,2) NOT NULL,
    S_COMMENT      VARCHAR(101) NOT NULL
);
```

```
CREATE TABLE PARTSUPP (
    PS_PARTKEY      INTEGER      NOT NULL,
    PS_SUPPKEY      INTEGER      NOT NULL,
    PS_AVAILQTY     INTEGER      NOT NULL,
    PS_SUPPLYCOST   DECIMAL(15,2) NOT NULL,
    PS_COMMENT      VARCHAR(199) NOT NULL
);
```

```
CREATE TABLE CUSTOMER (
    C_CUSTKEY      INTEGER      NOT NULL,
    C_NAME         VARCHAR(25)  NOT NULL,
    C_ADDRESS      VARCHAR(40)  NOT NULL,
    C_NATIONKEY    INTEGER      NOT NULL,
    C_PHONE        CHAR(15)     NOT NULL,
    C_ACCTBAL      DECIMAL(15,2) NOT NULL,
    C_MKTSEGMENT   CHAR(10)     NOT NULL,
    C_COMMENT      VARCHAR(117) NOT NULL
);
```

```
CREATE TABLE ORDERS (
    O_ORDERKEY      INTEGER      NOT NULL,
    O_CUSTKEY      INTEGER      NOT NULL,
    O_ORDERSTATUS   CHAR(1)     NOT NULL,
    O_TOTALPRICE    DECIMAL(15,2) NOT NULL,
    O_ORDERDATE     DATE        NOT NULL,
    O_ORDERPRIORITY CHAR(15)     NOT NULL,
    O_CLERK         CHAR(15)     NOT NULL,
    O_SHIPPRIORITY  INTEGER      NOT NULL,
    O_COMMENT      VARCHAR(79)  NOT NULL
);
```

```
);
```

```
-----

CREATE TABLE LINEITEM (
    L_ORDERKEY      INTEGER      NOT NULL,
    L_PARTKEY       INTEGER      NOT NULL,
    L_SUPPKEY       INTEGER      NOT NULL,
    L_LINENUMBER    INTEGER      NOT NULL,
    L_QUANTITY       DECIMAL(15,2) NOT NULL,
    L_EXTENDEDPRICE DECIMAL(15,2) NOT NULL,
    L_DISCOUNT     DECIMAL(15,2) NOT NULL,
    L_TAX           DECIMAL(15,2) NOT NULL,
    L_RETURNFLAG    CHAR(1)      NOT NULL,
    L_LINESTATUS    CHAR(1)      NOT NULL,
    L_SHIPDATE      DATE         NOT NULL,
    L_COMMITDATE    DATE         NOT NULL,
    L_RECEIPTDATE   DATE         NOT NULL,
    L_SHIPINSTRUCT  CHAR(25)     NOT NULL,
    L_SHIPMODE      CHAR(10)     NOT NULL,
    L_COMMENT       VARCHAR(44)  NOT NULL
);
```

Benchmark Queries

```
SELECT H.C_NAME, O.O_ORDERKEY, O.O_TOTALPRICE
FROM (
    SELECT CUSTOMER.C_CUSTKEY, CUSTOMER.C_NAME, CUSTOMER.C_ACCTBAL
    FROM CUSTOMER
    WHERE CUSTOMER.C_ACCTBAL > 1000
) H
JOIN (
    SELECT CUSTOMER.C_CUSTKEY, CUSTOMER.C_NAME, CUSTOMER.C_ACCTBAL
    FROM CUSTOMER
    WHERE CUSTOMER.C_ACCTBAL > 1000
) I ON H.CUSTOMER.C_NATIONKEY = I.CUSTOMER.C_NATIONKEY
JOIN ORDERS O ON H.CUSTOMER.C_CUSTKEY = O.O_CUSTKEY
WHERE O.O_TOTALPRICE > 50000 AND O.O_ORDERSTATUS = 'O' OR O.O_SHIPPRIORITY = 0;
```

```
-----

SELECT L.L_ORDERKEY, L.L_PARTKEY
FROM LINEITEM L
JOIN ORDERS O ON L.L_ORDERKEY = O.O_ORDERKEY
WHERE L.L_SHIPDATE >= '1994-01-01' AND O.O_ORDERSTATUS = 'F';
```

```
SELECT P.P_NAME, P.P_BRAND
FROM PART P
JOIN PARTSUPP PS ON P.P_PARTKEY = PS.PS_PARTKEY
JOIN SUPPLIER S ON PS.PS_SUPPKEY = S.S_SUPPKEY
WHERE P.P_SIZE >= 25;
```

```
SELECT S.S_NAME, N.N_NAME, L.L_EXTENDEDPRICE, O.O_ORDERDATE
FROM SUPPLIER S
JOIN NATION N ON S.S_NATIONKEY = N.N_NATIONKEY
JOIN LINEITEM L ON S.S_SUPPKEY = L.L_SUPPKEY
JOIN ORDERS O ON L.L_ORDERKEY = O.O_ORDERKEY
WHERE L.L_DISCOUNT > 0.05;
```

```
SELECT
    P.P_NAME,
    PS.PS_SUPPLYCOST,
    S.S_NAME,
    N.N_NAME,
    R.R_NAME
FROM PART P
JOIN PARTSUPP PS ON P.P_PARTKEY = PS.PS_PARTKEY
JOIN SUPPLIER S ON PS.PS_SUPPKEY = S.S_SUPPKEY
JOIN NATION N ON S.S_NATIONKEY = N.N_NATIONKEY
JOIN REGION R ON N.N_REGIONKEY = R.R_REGIONKEY
WHERE PS.PS_SUPPLYCOST < 300
    AND R.R_NAME = 'EUROPE';
```

```
SELECT P.P_NAME, S.S_NAME, N.N_NAME, PS.PS_SUPPLYCOST
FROM PART P
JOIN PARTSUPP PS ON P.P_PARTKEY = PS.PS_PARTKEY
JOIN SUPPLIER S ON PS.PS_SUPPKEY = S.S_SUPPKEY
JOIN NATION N ON S.S_NATIONKEY = N.N_NATIONKEY
WHERE P.P_TYPE = 'COPPER'
    AND PS.PS_SUPPLYCOST < 500;
```

```
SELECT H.S_NAME, I.S_PHONE, N.N_NAME
FROM (
    SELECT SUPPLIER.S_SUPPKEY, SUPPLIER.S_NAME, SUPPLIER.S_NATIONKEY
    FROM SUPPLIER
    WHERE SUPPLIER.S_ACCTBAL > 5000
) H
JOIN (
```

```
SELECT SUPPLIER.S_SUPPKEY, SUPPLIER.S_PHONE, SUPPLIER.S_NATIONKEY
FROM SUPPLIER
WHERE SUPPLIER.S_ACCTBAL > 5000
```

```
) I ON H.SUPPLIER.S_NATIONKEY = I.SUPPLIER.S_NATIONKEY
JOIN NATION N ON H.SUPPLIER.S_NATIONKEY = N.N_NATIONKEY
WHERE N.N_NAME = 'UNITED STATES';
-----
```

```
SELECT O.ORDERKEY, L.QUANTITY, S.SUPPKEY, C.CUSTKEY
FROM ORDERS O
JOIN CUSTOMER C ON O.CUSTKEY = C.CUSTKEY
JOIN LINEITEM L ON O.ORDERKEY = L.ORDERKEY
JOIN SUPPLIER S ON S.SUPPKEY = L.SUPPKEY
JOIN (SELECT N.NATIONKEY, N.NAME FROM NATION N) TMP1 ON TMP1.NATIONKEY = C.NATIONKEY
JOIN (SELECT N.NATIONKEY, N.NAME FROM NATION N) TMP2 ON TMP2.NATIONKEY = S.NATIONKEY;
-----
```

```
SELECT PS.PARTKEY, PS.SUPPLYKEY
FROM PARTSUPP PS
JOIN SUPPLIER S ON PS.SUPPLYKEY = S.SUPPLYKEY
JOIN LINEITEM L ON PS.SUPPLYKEY = L.SUPPLYKEY
JOIN (SELECT P.NAME, P.BRAND FROM PART P) TMP1 ON TMP1.PARTKEY = PS.PARTKEY
JOIN (SELECT P.NAME, P.BRAND FROM PART P) TMP2 ON TMP2.PARTKEY = L.PARTKEY
WHERE PS.AVAILQTY > 10 AND S.ACCTBAL > 1000;
-----
```

```
SELECT P.PARTKEY, P.SUPPKEY
FROM PARTSUPP P
JOIN SUPPLIER S ON P.SUPPKEY = S.SUPPKEY
WHERE P.AVAILQTY > 10 OR P.SUPPLYCOST < 500 AND S.ACCTBAL > 1000;
-----
```